

It Takes Two to Tango: Mixup for Deep Metric Learning



Shashanka
Venkataramanan



Bill
Psomas



Ewa
Kijak



Laurent
Amsaleg



Konstantinos
Karantzalos




Yannis
Avrithis



Deep Metric Learning


- GOAL – Learning a **discriminative representation** that generalizes to unseen classes.
- HOW? – Intra-class embeddings are **pulled** closer and inter-class embeddings are **pushed** apart.
- MOTIVATION – Classes during training and inference are **different**, interpolation-based data augmentation e.g. mixup plays **significant** role.



$$L_{cont} = \sum_{p \in P(a)} -s(a, p) + \sum_{n \in N(a)} [s(a, n) - m]_+$$

Interpolation for pairwise loss functions


Embedding Expansion



Interpolate pairs of embeddings in a deterministic way **within the same class**.

Do not perform **label interpolation**.


Proxy Synthesis



Interpolates **between classes**, applying to proxy-based losses only.

risks synthesizing false negatives when the interpolation factor λ is close to 0 or 1.

MoCHI



Interpolates **anchor with negative** embeddings.

do not interpolate labels, chooses $\lambda \in [0, 0.5]$ to avoid false negatives.

what is a proper way to **define** and **interpolate** labels
for deep metric learning ?


what is a proper way to **define** and **interpolate** labels
for deep metric learning ?



Metrix

Improving Representations Using Mixup

- Mixup: a data augmentation technique that **interpolates between two examples** (input or feature) and its **corresponding labels**.



$$\lambda \sim \text{Beta}(\alpha, \alpha)$$
$$\text{mix}_{\lambda}(a, b) = \lambda a + (1 - \lambda)b'$$

Generic Loss Formulation

- Additive losses e.g., Contrastive and non-additive e.g., Multi-similarity involve a **sum over positives** $P(a)$ and a **sum over negatives** $N(a)$.

$$\ell(a; \theta) := \tau \left(\sigma^+ \left(\sum_{p \in P(a)} \rho^+(s(a, p)) \right) + \sigma^- \left(\sum_{n \in N(a)} \rho^-(s(a, n)) \right) \right)$$

sum over positives sum over negatives

Table 1 of our paper, shows the values of each of these terms for different loss functions.

Generic Loss Formulation

- Additive losses e.g., Contrastive and non-additive e.g., Multi-similarity involve a **sum over positives** $P(a)$ and a **sum over negatives** $N(a)$.
- They also involve a **decreasing function** of similarity $s(a, p) \forall p \in P(a)$ and an increasing function of similarity $s(a, n) \forall n \in N(a)$.

$$\ell(a; \theta) := \tau \left(\sigma^+ \left(\sum_{p \in P(a)} \boxed{\rho^+}(s(a, p)) \right) + \sigma^- \left(\sum_{n \in N(a)} \boxed{\rho^-}(s(a, n)) \right) \right)$$

decreasing function of similarity increasing function of similarity

Generic Loss Formulation

- Additive losses e.g., Contrastive and non-additive e.g., Multi-similarity involve a **sum over positives** $P(a)$ and a **sum over negatives** $N(a)$.
- They also involve a **decreasing function** of similarity $s(a, p) \forall p \in P(a)$ and an increasing function of similarity $s(a, n) \forall n \in N(a)$.

$$\ell(a; \theta) := \tau \left(\sigma^+ \left(\sum_{p \in P(a)} \rho^+(s(a, p)) \right) + \sigma^- \left(\sum_{n \in N(a)} \rho^-(s(a, n)) \right) \right)$$

non-linear functions
(for non-additive losses)

Table 1 of our paper, shows the values of each of these terms for different loss functions.

Generic Loss Formulation

- In metric learning, positives $P(a)$ and negatives $N(a)$ of anchor a have the **same** or **different** class label as the anchor.
- We assign binary class label $y \in \{0,1\} \forall P(a) \cup N(a)$ s.t. **$y = 1$ for positives** and **$y = 0$ for negatives**.

$$\ell(a; \theta) := \tau \left(\sigma^+ \left(\sum_{(x,y) \in U(a)} \boxed{y} \rho^+(s(a, x)) \right) + \sigma^- \left(\sum_{(x,y) \in U(a)} \boxed{(1-y)} \rho^-(s(a, x)) \right) \right)$$

y is binary, **only one of the two** contributions is nonzero.

Interpolating Labels Using Generic Formulation

- Given $M(a)$ which is the possible choices of mixing pairs (positive-positive or positive-negative or negative-negative), the labeled mixed embeddings is


$$V(a) = \{f_\lambda(x, x'), \text{mix}_\lambda(y, y') : ((x, y), (x', y')) \in M(a)\}$$

$$\tilde{\ell}(a; \theta) := \tau \left(\sigma^+ \left(\sum_{(v, y) \in V(a)} \boxed{y} \rho^+(s(a, v)) \right) + \sigma^- \left(\sum_{(v, y) \in V(a)} \boxed{(1 - y)} \rho^-(s(a, v)) \right) \right)$$


$y \in [0, 1]$, both contributions are nonzero.

Comparison With Other Mixing Methods

Mixup in input space

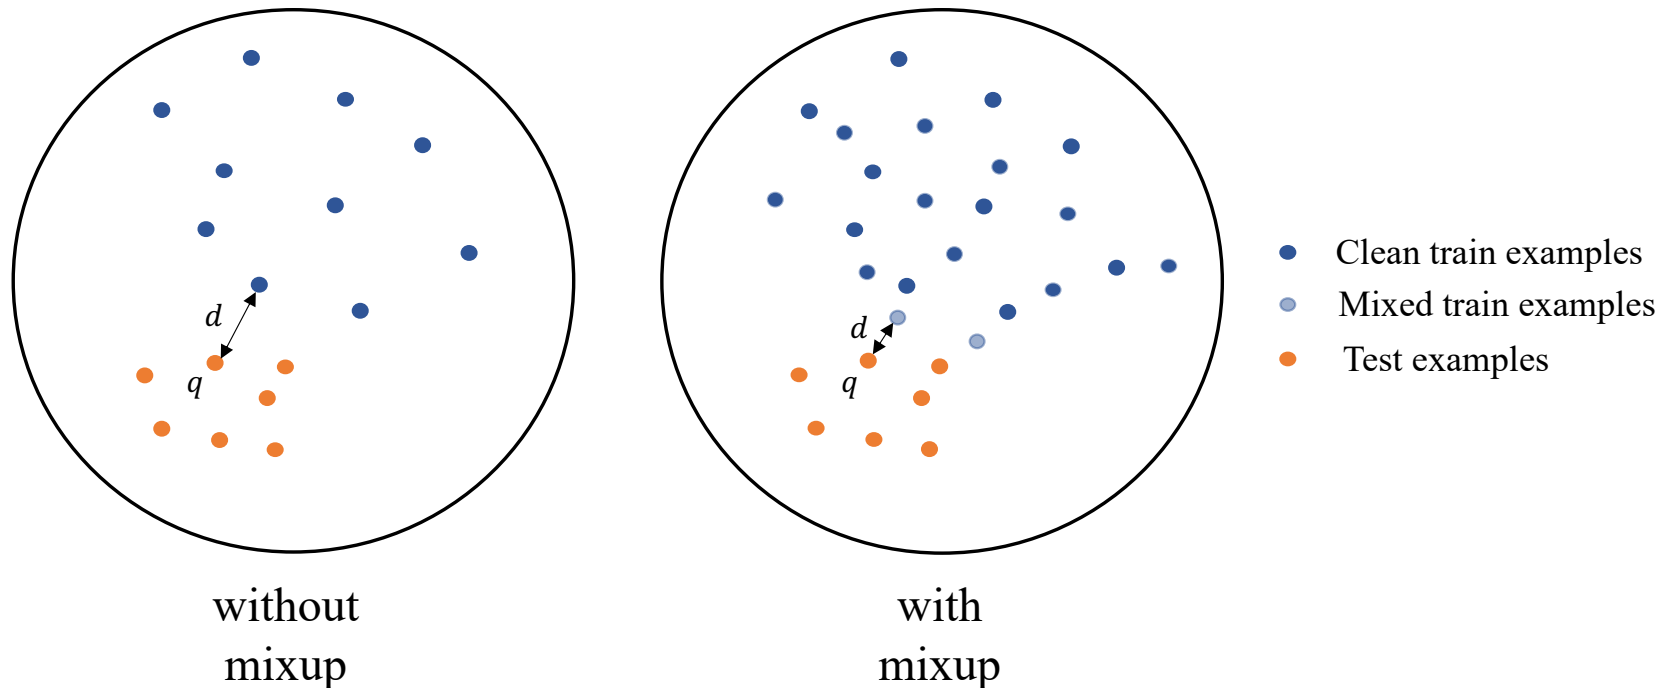


Mixup in embedding space



How Does Mixup Improve Representations?

- Introduce a new evaluation metric - **utilization** and show that a representation more appropriate for test classes is implicitly learned during **exploration of the embedding space** in the presence of mixup.





Paper



Code